

# CS 241 Honors Memory

Ben Kurtovic   Atul Sandur   Bhuvan Venkatesh   Brian Zhou  
Kevin Hong

University of Illinois Urbana–Champaign

February 21, 2017

- Memory review
- Management techniques
  - Heap vs. stack
  - `free()`
  - `alloca()`
  - Smart pointers (C++)
- Garbage collection
  - Reference counting
  - Mark-and-sweep
  - Copying collector
  - Generational
  - In practice
- Concluding thoughts

# Types of variables in program memory

- Static variables
  - Persist throughout the lifetime of the program
  - No runtime cost, but memory cannot be used for anything else

# Types of variables in program memory

- Static variables
  - Persist throughout the lifetime of the program
  - No runtime cost, but memory cannot be used for anything else
- Local variables, function arguments
  - Live for the lifetime of the function call
  - Stored on the stack
  - Allocation and reclaiming of memory is really cheap—why?

# Types of variables in program memory

- Static variables
  - Persist throughout the lifetime of the program
  - No runtime cost, but memory cannot be used for anything else
- Local variables, function arguments
  - Live for the lifetime of the function call
  - Stored on the stack
  - Allocation and reclaiming of memory is really cheap—why?
- Heap allocation
  - Slower than stack allocation because of bookkeeping by memory manager
  - This memory is not automatically reclaimed, so will eventually fill up
  - Call `free()`! But it isn't always obvious when you are done with something

# malloc() and free()

What's the big deal?

You all know from using it—manual heap memory management is a pain

# Why is heap management so difficult?

- Manual bookkeeping (programming error)
  - Dangling pointers, double deletion, memory leaks
  - Don't always know when to clean up memory

# Why is heap management so difficult?

- Manual bookkeeping (programming error)
  - Dangling pointers, double deletion, memory leaks
  - Don't always know when to clean up memory
- Performance tradeoffs with heap memory management
  - Linked list of free blocks, how expensive to get a new block of memory to allocate?
  - Pre-coalesce blocks so allocation can be faster? free becomes expensive



# Why is heap management so difficult?

- Manual bookkeeping (programming error)
  - Dangling pointers, double deletion, memory leaks
  - Don't always know when to clean up memory
- Performance tradeoffs with heap memory management
  - Linked list of free blocks, how expensive to get a new block of memory to allocate?
  - Pre-coalesce blocks so allocation can be faster? free becomes expensive
- Fragmentation
  - Bad locality of reference
  - Problem in virtual memory systems due to page faults, cache misses

# free(), continued

Keeping mental track of your heap memory gets frustrating in complex programs, especially if you're trying to be careful about error handling

```
a = malloc(...);
b = malloc(...);

if (/* do something and it fails */) {
    free(a);
    free(b);
    return;
}

if (/* do another thing */) {
    if (/* do a third thing and it fails */) {
        free(a);
        free(b);
        return;
    }
    /* ... */
}

c = malloc(...);
if (/* do a fourth thing and it fails */) {
    free(c);
    free(b);
    free(a);
    return;
}

/* ... */
```

# free(), continued (the "goto cleanup" idiom)

```
a = malloc(...);
b = malloc(...);

if (/* do something and it fails */)
    goto cleanup;

if (/* do another thing */) {
    if (/* do a third thing and it fails */)
        goto cleanup;
    /* ... */
}

c = malloc(...);
if (/* do a fourth thing and it fails */)
    goto cleanup;

/* ... */

cleanup:
free(a);
free(b);
free(c);
```

# free(), continued (the "goto cleanup" idiom)

```
a = malloc(...);
b = malloc(...);

if (/* do something and it fails */)
    goto cleanup;

if (/* do another thing */) {
    if (/* do a third thing and it fails */)
        goto cleanup;
    /* ... */
}

c = malloc(...);
if (/* do a fourth thing and it fails */)
    goto cleanup;

/* ... */

cleanup:
free(a);
free(b);
free(c);
```

- Reduces usage of free() and makes keeping track of variables easier—only need to think about it at the end of the function
- Can be generalized to other resource types (opening/closing files)

# free(), continued (the "goto cleanup" idiom)

```
a = malloc(...);
b = malloc(...);

if (/* do something and it fails */)
    goto cleanup;

if (/* do another thing */) {
    if (/* do a third thing and it fails */)
        goto cleanup;
    /* ... */
}

c = malloc(...);
if (/* do a fourth thing and it fails */)
    goto cleanup;

/* ... */

cleanup:
free(a);
free(b);
free(c);
```

- Reduces usage of free() and makes keeping track of variables easier—only need to think about it at the end of the function
- Can be generalized to other resource types (opening/closing files)
- Makes things easier, but still have manual heap management

# So, how can we take advantage of the stack?

# So, how can we take advantage of the stack?

Let's take a look at some of its limitations...

# alloca()

Stack memory is cheap, but often limited by fixed sizes

```
char buffer[512];  
puts("What's your name?");  
fgets(buffer, sizeof(buffer), stdin);
```



# alloca()

Stack memory is cheap, but often limited by fixed sizes

```
char buffer[512];  
puts("What's your name?");  
fgets(buffer, sizeof(buffer), stdin);
```

What if the person has a **really** long name?

# alloca()

Stack memory is cheap, but often limited by fixed sizes

```
char buffer[512];  
puts("What's your name?");  
fgets(buffer, sizeof(buffer), stdin);
```

What if the person has a **really** long name?

```
if (/* entire name wasn't read */) {  
    buffer = realloc(buffer, 2 * sizeof(buffer));  
    fgets(...);  
}
```

# alloca()

Stack memory is cheap, but often limited by fixed sizes

```
char buffer[512];  
puts("What's your name?");  
fgets(buffer, sizeof(buffer), stdin);
```

What if the person has a **really** long name?

```
if (/* entire name wasn't read */) {  
    buffer = realloc(buffer, 2 * sizeof(buffer));  
    fgets(...);  
}
```

```
*** Error in './name': realloc(): invalid pointer: 0x00007ffffb3af9250 ***  
===== Backtrace: =====  
/lib64/libc.so.6(+0x7bc17)[0x7f02fc21dc17]  
/lib64/libc.so.6(realloc+0x31e)[0x7f02fc2228fe]  
...  
Aborted (core dumped)
```

You can't realloc() stack memory

# alloca()

Stack memory is cheap, but often limited by fixed sizes

```
char buffer[512];  
puts("What's your name?");  
fgets(buffer, sizeof(buffer), stdin);
```

What if the person has a **really** long name?

```
if (/* entire name wasn't read */) {  
    buffer = realloc(buffer, 2 * sizeof(buffer));  
    fgets(...);  
}
```

```
*** Error in './name': realloc(): invalid pointer: 0x00007ffffb3af9250 ***  
===== Backtrace: =====  
/lib64/libc.so.6(+0x7bc17)[0x7f02fc21dc17]  
/lib64/libc.so.6(realloc+0x31e)[0x7f02fc2228fe]  
...  
Aborted (core dumped)
```

You can't `realloc()` stack memory... *or can you?*

# alloca()

Stack memory is cheap, but often limited by fixed sizes

```
char buffer[512];  
puts("What's your name?");  
fgets(buffer, sizeof(buffer), stdin);
```

What if the person has a **really** long name?

```
if (/* entire name wasn't read */) {  
    buffer = realloc(buffer, 2 * sizeof(buffer));  
    fgets(...);  
}
```

```
*** Error in './name': realloc(): invalid pointer: 0x00007ffffb3af9250 ***  
===== Backtrace: =====  
/lib64/libc.so.6(+0x7bc17)[0x7f02fc21dc17]  
/lib64/libc.so.6(realloc+0x31e)[0x7f02fc2228fe]  
...  
Aborted (core dumped)
```

You can't `realloc()` stack memory... *or can you?* (you can't)

# alloca()

But, there *is* another way

```
#include <alloca.h>

void *alloca(size_t size);
```

# alloca()

But, there *is* another way

```
#include <alloca.h>

void *alloca(size_t size);
```

Use it exactly like `malloc()`, but memory is allocated on the stack rather than heap

- No `free()` necessary! (It would crash, like before)
- Thus, memory leaks are impossible
- Cheap and fast (no heap memory management)

# alloca()

But, there *is* another way

```
#include <alloca.h>

void *alloca(size_t size);
```

Use it exactly like `malloc()`, but memory is allocated on the stack rather than heap

- No `free()` necessary! (It would crash, like before)
- Thus, memory leaks are impossible
- Cheap and fast (no heap memory management)

What's wrong with this?



# alloca()

But, there *is* another way

```
#include <alloca.h>

void *alloca(size_t size);
```

Use it exactly like `malloc()`, but memory is allocated on the stack rather than heap

- No `free()` necessary! (It would crash, like before)
- Thus, memory leaks are impossible
- Cheap and fast (no heap memory management)

What's wrong with this?

- Stack scoping rules: can't return a pointer to stack space!
- Overflow: stack space is limited; can't allocate too much memory

# alloca()

But, there *is* another way

```
#include <alloca.h>

void *alloca(size_t size);
```

Use it exactly like `malloc()`, but memory is allocated on the stack rather than heap

- No `free()` necessary! (It would crash, like before)
- Thus, memory leaks are impossible
- Cheap and fast (no heap memory management)

What's wrong with this?

- Stack scoping rules: can't return a pointer to stack space!
- Overflow: stack space is limited; can't allocate too much memory
- Non-standard! (Not part of ANSI C, POSIX, etc...)

# alloca()

But, there *is* another way

```
#include <alloca.h>

void *alloca(size_t size);
```

Use it exactly like malloc(), but memory is allocated on the stack rather than heap

- No free() necessary! (It would crash, like before)
- Thus, memory leaks are impossible
- Cheap and fast (no heap memory management)

What's wrong with this?

- Stack scoping rules: can't return a pointer to stack space!
- Overflow: stack space is limited; can't allocate too much memory
- Non-standard! (Not part of ANSI C, POSIX, etc...)

So, not a good idea

# C++ smart pointers

New feature in C++11 standard library (although concept existed before)

# C++ smart pointers

New feature in C++11 standard library (although concept existed before)

Concept of *ownership*:

# C++ smart pointers

New feature in C++11 standard library (although concept existed before)

Concept of *ownership*:

- Instead of passing pointers around like mad, let's set some rules

New feature in C++11 standard library (although concept existed before)

Concept of *ownership*:

- Instead of passing pointers around like mad, let's set some rules
- RAI (Resource Acquisition Is Initialization)—remember from CS 225, maybe?
  - Only allocate memory in your constructor and free it in your destructor
  - Put things on the stack as much as possible

New feature in C++11 standard library (although concept existed before)

Concept of *ownership*:

- Instead of passing pointers around like mad, let's set some rules
- RAI (Resource Acquisition Is Initialization)—remember from CS 225, maybe?
  - Only allocate memory in your constructor and free it in your destructor
  - Put things on the stack as much as possible
- "Smart pointer" wraps an ordinary (raw) pointer and defines the semantics for how it is managed and who's managing it
- Using these rules, you should (almost) never need to manually `new/delete` memory



## Types of smart pointers

- `std::unique_ptr`: Exactly one owner (compile-time error to try to copy it)

When that owner lets the pointer go out of scope, the memory is automatically released

## Types of smart pointers

- `std::unique_ptr`: Exactly one owner (compile-time error to try to copy it)  
When that owner lets the pointer go out of scope, the memory is automatically released
- `std::shared_ptr`: Reference-counted version (we'll discuss what this means later)  
In other words, it can have multiple owners, and is released when the last one lets go

## Types of smart pointers

- `std::unique_ptr`: Exactly one owner (compile-time error to try to copy it)  
When that owner lets the pointer go out of scope, the memory is automatically released
- `std::shared_ptr`: Reference-counted version (we'll discuss what this means later)  
In other words, it can have multiple owners, and is released when the last one lets go
- `std::weak_ptr`: Stores a "temporary" reference without owning it  
Can be used to break cycles if two objects refer to each other (e.g., doubly linked lists)  
Useful for caches, which store references to objects that *should* be deleted if the cache is the only thing still referring to it

# C++ smart pointers

Old way:

```
bool func() {
    size_t bufsize = 1024 * 1024 * 128;
    int* data = new int[bufsize];

    for (int i = 0; i < bufsize; i++) {
        data[i] = get_data();
        if (!data[i]) {
            delete[] data;
            return false;
        }
    }

    // do something fancy with data
    delete[] data;
    return true;
}
```

# C++ smart pointers

## Old way:

```
bool func() {
    size_t bufsize = 1024 * 1024 * 128;
    int* data = new int[bufsize];

    for (int i = 0; i < bufsize; i++) {
        data[i] = get_data();
        if (!data[i]) {
            delete[] data;
            return false;
        }
    }

    // do something fancy with data
    delete[] data;
    return true;
}
```

## Using std::unique\_ptr:

```
bool func() {
    size_t bufsize = 1024 * 1024 * 128;
    auto data = std::make_unique<int[]>(
        bufsize);

    for (int i = 0; i < bufsize; i++) {
        data[i] = get_data();
        if (!data[i])
            return false;
    }

    // do something fancy with data
    return true;
}
```

# C++ smart pointers

Old way:

```
bool func() {
    size_t bufsize = 1024 * 1024 * 128;
    int* data = new int[bufsize];

    for (int i = 0; i < bufsize; i++) {
        data[i] = get_data();
        if (!data[i]) {
            delete[] data;
            return false;
        }
    }

    // do something fancy with data
    delete[] data;
    return true;
}
```

Using `std::unique_ptr`:

```
bool func() {
    size_t bufsize = 1024 * 1024 * 128;
    auto data = std::make_unique<int[]>(
        bufsize);

    for (int i = 0; i < bufsize; i++) {
        data[i] = get_data();
        if (!data[i])
            return false;
    }

    // do something fancy with data
    return true;
}
```

No calls to `new[]` or `delete[]` necessary

C++ will automatically allocate the memory we want when we call `std::make_unique`, and release it when the unique pointer goes out of scope when the function exits

...but not everyone uses C++

# What about when you *need* the heap?

- You can't always rely on the stack because it can overflow quickly
- How do you easily manage large objects?



# What about when you *need* the heap?

- You can't always rely on the stack because it can overflow quickly
- How do you easily manage large objects?
- Garbage collection to the rescue! - Automagic Memory Management

# GC: Reference counting

- Each object has count of how many objects point to it
- Need to track count whenever object pointers are set or removed/freed

## Manual Reference Counting



```
MyClass *obj1 = [[MyClass alloc] init];
```



```
MyClass *obj2 = [obj1 retain];
```



```
[obj2 release];
```



```
[obj1 release];
```

Python!

# GC: Reference Counting Problems


- Overhead of maintaining count

# GC: Reference Counting Problems

- Overhead of maintaining count
- Lumpy deletions: tend to have clusters of related objects.

# GC: Reference Counting Problems

- Overhead of maintaining count
- Lumpy deletions: tend to have clusters of related objects.
- Cycles



- Most common form of garbage collection

- Most common form of garbage collection
- "Trace" connections between references and allocations



- Most common form of garbage collection
- "Trace" connections between references and allocations

# GC: Mark-and-Sweep

- Two Step Process

# GC: Mark-and-Sweep

- Two Step Process
- Mark - Find and label all accessible objects

# GC: Mark-and-Sweep

- Two Step Process
- Mark - Find and label all accessible objects
  - Follow all pointers from the "root set".
  - Perform DFS on pointers
  - Flag field set to marked

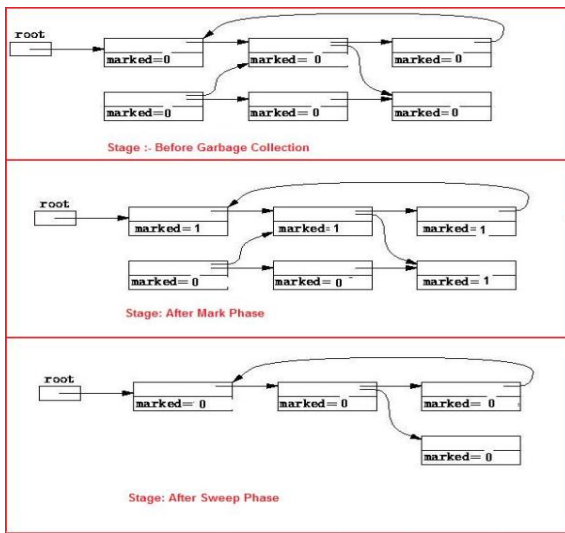
# GC: Mark-and-Sweep

- Sweep - Remove all inaccessible objects

# GC: Mark-and-Sweep

- Sweep - Remove all inaccessible objects
  - Iterate through memory
  - If marked, unmark
  - Else, free

# GC: Mark-and-Sweep



# GC: Mark-and-Sweep

- Pros



# GC: Mark-and-Sweep

- Pros
  - Works with cyclic data structures
  - Effective and easy to implement

# GC: Mark-and-Sweep

- Pros
  - Works with cyclic data structures
  - Effective and easy to implement
- Cons

# GC: Mark-and-Sweep


- Pros
  - Works with cyclic data structures
  - Effective and easy to implement
- Cons
  - Program must halt
  - Fragmentation

# GC: Mark-and-Sweep with Compaction

- Alleviates fragmentation
- Sweep moves memory to the left

# GC: Mark-and-Sweep with Compaction

- Alleviates fragmentation
- Sweep moves memory to the left



# GC: Tri-Color Marking

- Incremental garbage collector

# GC: Tri-Color Marking

- Incremental garbage collector
- Gray Set - Need to be scanned blocks
- Black Set - Accessible blocks
- White Set - Inaccessible blocks


# GC: Tri-Color Marking

- Incremental garbage collector
- Gray Set - Need to be scanned blocks
- Black Set - Accessible blocks
- White Set - Inaccessible blocks
- All blocks accessible from the root start in the gray set.
- All others start in the white set.
- Gray moves to black, all white set objects referenced moved to gray.



# Generational GC

- Do we really need to scan everything all the time?
- Split the objects up into generations
- Scan each generation with a different frequency



# Generational GC - Eden

- Compaction not important, needs fast GC
- When full, run mark phase and move objects up



# Generational GC - Survivor Space 1

- Count for objects being copied back and forth
- When full, run mark phase and move objects up



# Generational GC - Survivor Space 2

- Compaction not important
- Less frequent GC



# Generational GC - Old Generation

- Mark/sweep with compaction since objects live longer
- Run less frequently, since slow process



# Generational GC - Problems

- Objects in old generation could be pointing to objects in survivor space
- Mark phase in survivor space does not go through objects in old generation
- Maintain separate table for maintaining such pointers across generations
- Turns out such occurrences are less frequent in code bases
- Objects in old generation don't generally change much, so may not be too expensive in practice
- Add references from this table to root set of copying GC running in survivor space
- What about younger→older generation references?

# Garbage collection in C/C++?

- Doesn't come with one for performance reasons!
- Usually a bolted-on third party library
- Link the library and allow it to reclaim unused memory automatically
- `malloc` can be replaced with garbage collector's allocator, and `free` is a no-op

# Garbage collection in C/C++?

```
#include <assert.h>
#include <stdio.h>
#include <gc.h>

int main(void)
{
    int i;
    const size = 100000000;

    GC_INIT();
    for (i = 0; i < size; ++i)
    {
        int **p = GC_MALLOC(sizeof *p);
        int *q = GC_MALLOC_ATOMIC(sizeof *q);

        assert(*p == 0);
        *p = GC_REALLOC(q, 2 * sizeof *p);
        if (i == size-1)
            printf("Heap size = %zu\n", GC_get_heap_size());
    }

    return 0;
}
```



# GC doesn't mean you forget about memory

- Explicitly let go of object by setting the reference to NULL. This is useful with global variables. Otherwise you hang onto system resources longer.
- Need to manage other system resources (like files); they are closed before GC by programmer

# Conclusion

Thank you! Questions?